Microsoft IQ: why context wins the AI deal

Knowit, Nordics · 28 October 2026

Johan Wallquist

Partner Tech Strategist, Microsoft

Ben Gibbs

Microsoft Global Black Belt

Another week. Another model.

What changed for your customer?



In four years, AI went from talking to doing

Nov 2022	**Chat**	One RFP question. One paragraph.
Late 2024	**Reasoning**	Analyze requirements. Develop a response strategy.
Late 2025	**Agents**	Research, draft, validate, and assemble the proposal.
2026+	**Agentic enterprise**	Orchestrate the entire RFP process across people, agents, and enterprise systems.

Chat answers. An agent plans, acts and checks.



Same model. Different system around it.

78%

Claude Code harness

42%

CORE-Agent harness

Claude Opus 4.5 · CORE-Bench

Agents reproduce published scientific results by running code.

Not an RFP benchmark or Microsoft IQ test. Different harnesses; the gap is not context alone.

Your role changes as AI does more of the work

You write it	AI suggests a sentence.
You fix it	AI drafts. You edit.
You brief it	A bounded pre-study, ready for review.
You run it	Parallel work. You set the rules.

Software and consulting need different context

Software

A repository gives a working context

Tests can check many outputs

Consulting

Context crosses teams and systems

Capacity or margin can mean different things

When AI drafts and you edit, you still fill the context gaps.



Capable models are widely available.

What does this customer know that the model doesn't?



Onboarding

A new colleague gets

Laptop and access

Who's who

The project history

A buddy

An agent also needs

Relevant context

Clear access and task boundaries

Microsoft IQ

WHAT KNOWIT CAN DELIVER

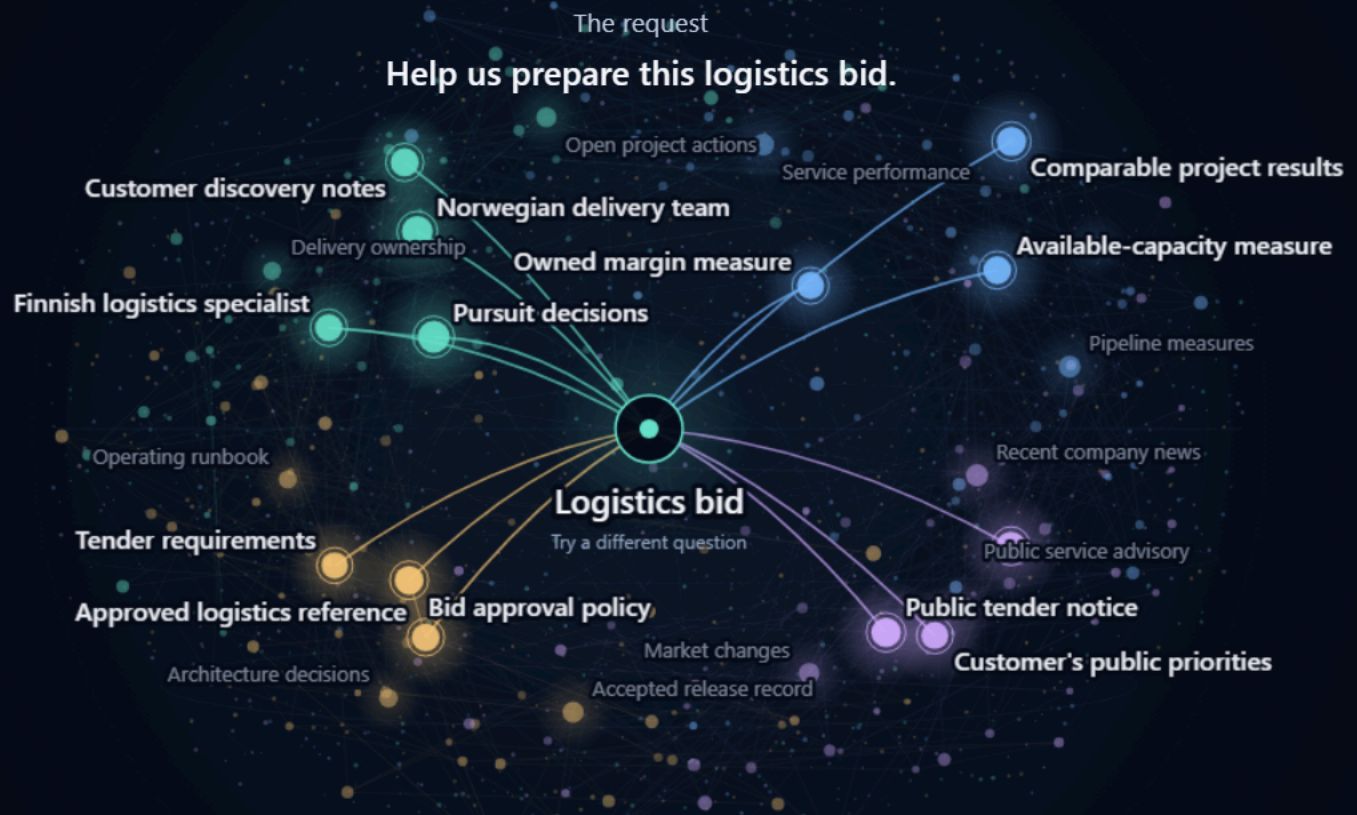
Where else could Knowit use this?

A handover needs different information from a bid. Reuse the sources that fit, with the right access.

RFP responses	Meeting preparation
Finding expertise	Project handovers
Service cases	Leadership decisions

A paid customer engagement

Connect the sources and agree the business definitions. Build the workflow and keep it working.



Change the question. A different set of information matters.

Use the capabilities the task needs. All four are not required for every project.

Work IQ	Fabric IQ	Foundry IQ	Web IQ
People & work	Business meaning & measures	Reusable chosen knowledge	Public web

Fictional Knowit example. Illustrative relationships and prepared outputs. Configured sources keep their own access controls. No live data is queried.

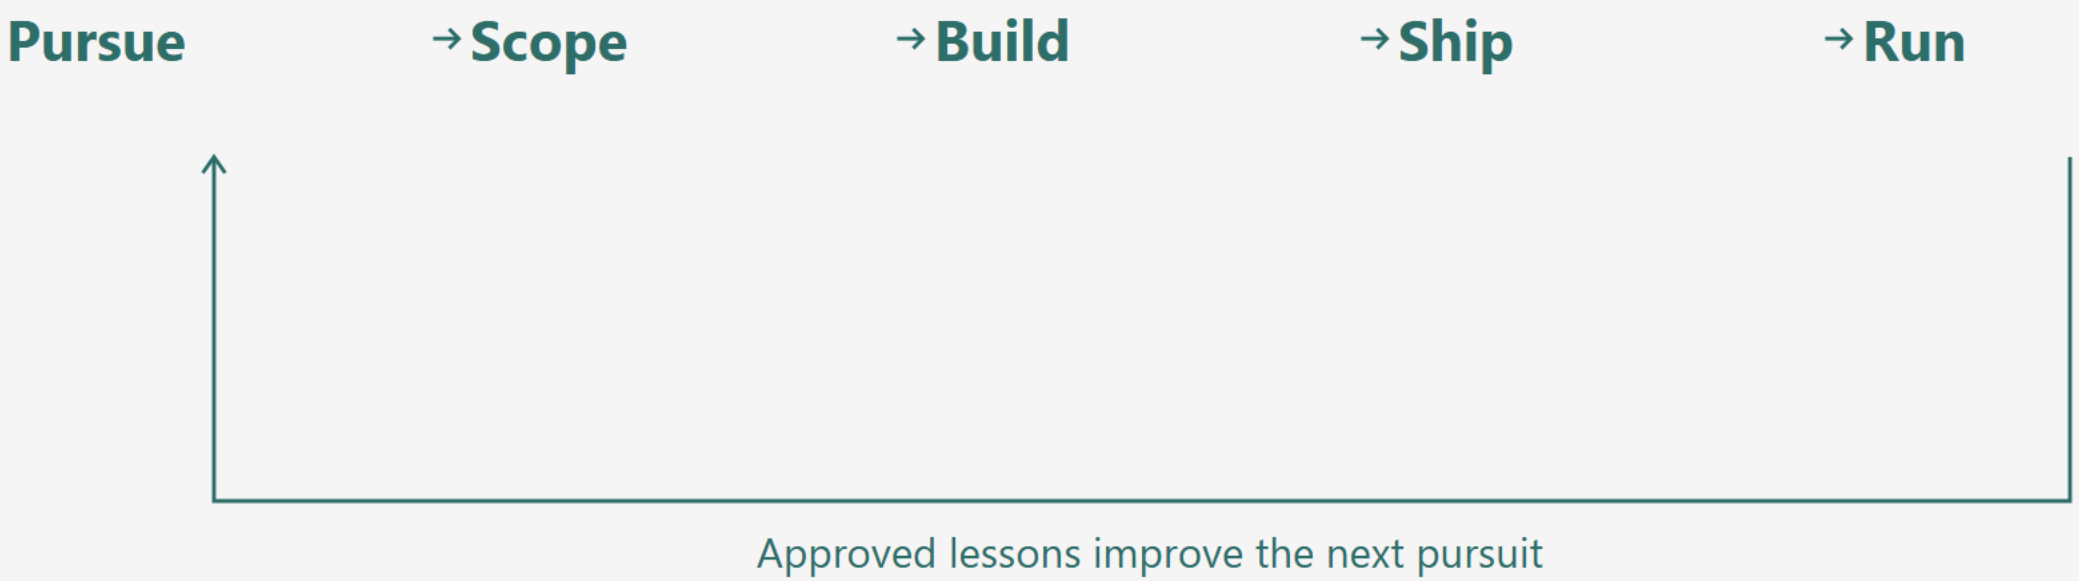
Check what's available before you promise it

Scope	The exact capability and experience
Availability	GA, preview or restricted access
Commercial terms	Licensing, usage and source connections

Verify the exact capability and terms before a customer proposal.

Knowit as customer zero

Try one repeated task. Measure useful work.



Proposed learning motion. No Knowit deployment or pilot is claimed.

One Knowit: combine what we already know

A fictional logistics pursuit

Sweden

A logistics lead

Norway

Approved delivery experience

Finland

Relevant expertise

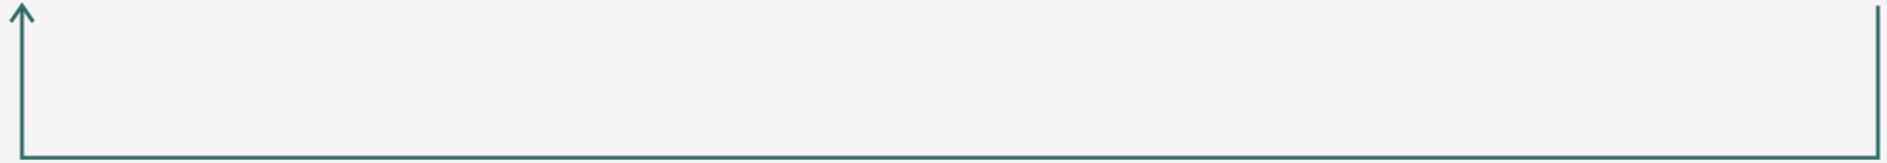
Data · Workplace · Applications

One brief, with an owner

Fictional pursuit. Sources and connections are illustrative.

Keep the context when the team changes

Pursue	→ Scope	→ Build	→ Ship	→ Run
Approved evidence and promises	Agreed outcomes and boundaries	Decisions and checks	Accepted release and ownership	Actual results improve the next pursuit



Approved lessons improve the next pursuit

Same fictional engagement. An owner at each stage.

Watch what the context does

Evidence	What context supports the answer?
Decisions	Where does a person approve?
Economics	What does it cost to run?

Over to Ben Gibbs.



Chief-of-Staff agent for sales

Ben Gibbs

Ben’s live Chief-of-Staff section

Shared context for sales leadership

RFP agent for sales

Ben Gibbs

Ben’s live RFP section

Preparing a bid response

What does it cost to run? How do you bring it down?

Ben Gibbs

Ben’s live AI FinOps section

Customer agent-running costs and evaluations

Where should Knowit look for opportunities?

Bids chase approved answers	RFP responses from agreed evidence.
Handovers rebuild decisions	Reusable scope and delivery context.
AI answers disagree	Shared definitions, sources and access.

Ask who owns the workflow, which sources count and how an approved output is measured.

Why build on Microsoft?

Already using Microsoft? Less context to rebuild.

Work IQ

Reuse Microsoft 365 work context.

Fabric / Power BI

Reuse definitions and trusted measures.

Foundry IQ

Reuse chosen knowledge with citations.

Configured controls

Identity and source access.

Reuse needs maintained context and configured access. Keep other platforms where they fit.

What can Knowit sell?

Workflow discovery	Baseline, source/access map and agreed definitions.
Workflow delivery	Cited RFP or handover drafts, human approval and acceptance tests.
Evaluation service	Maintained sources; quality/cost reviews against accepted outputs.

Proposed offers to shape with Knowit. Agree scope, fee and consumption explicitly.

Start with an existing Microsoft account

In the chat: customer or project + why

First conversation	One workflow + owner + sources
Bring	Knowit account lead + Data / M365 specialists
First engagement	Scope one offer. Use measured results to expand or reuse.

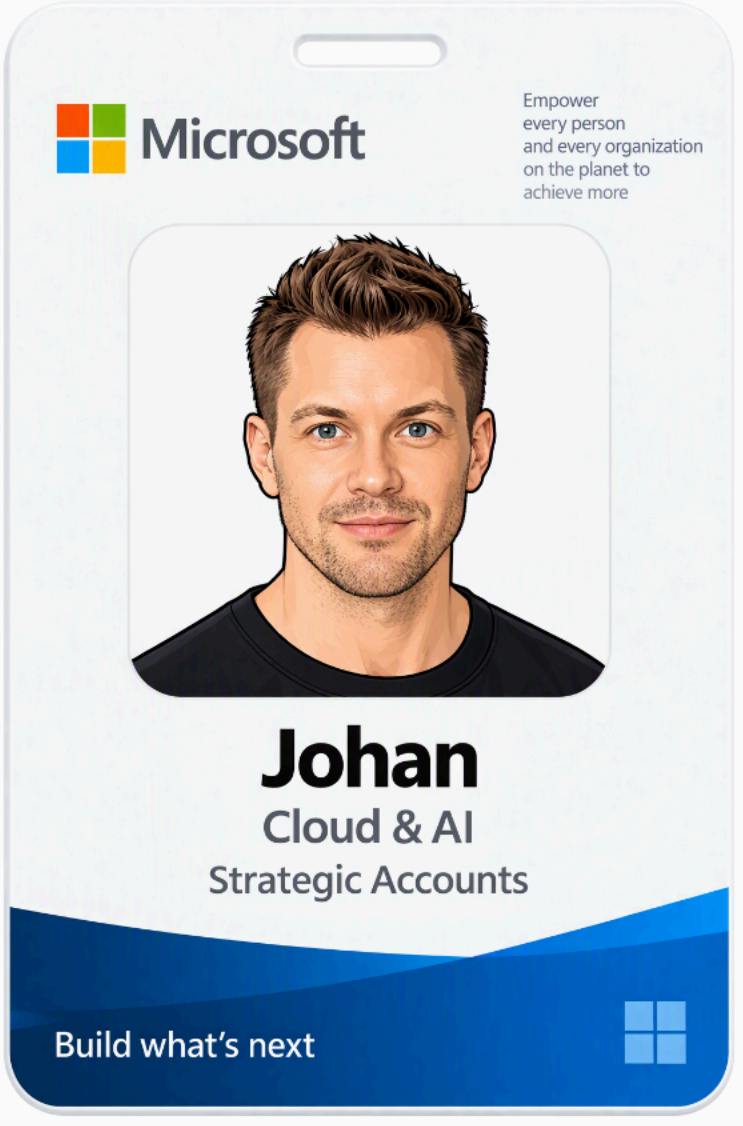
A safe description or private message is fine. No confidential detail.

Thank you

Connect and get the slides

bit.ly/jw-ai

LinkedIn, contact and resources



Johan Wallquist

Partner Tech Strategist, Microsoft